

Argus

An AI-assisted trading agent that sells insurance on the S&P 500, and publishes the honest math on whether it works.

LIVE ACCOUNT, REAL TRADES

SOURCE: WRITEUP.MD, OPENING LINE



SECTION 01

The Problem

Why options income is easy to sell and hard to do safely, and
why nobody checks the safety.

THE PROBLEM

Nobody tests the safety rails

Selling options for income is a proven strategy. Not blowing up doing it is the part that actually fails.

- By hand it takes discipline few people sustain: one oversized trade after a good month erases a year of income
- Handed to a bot, you're trusting risk rules nobody ever measured against real market history
- An untested safety rule is a guess wearing a safety label, and you find out when it costs you money

TYPICAL BOT'S RISK CONTROLS

~~Trust us~~


...tested against what?

THE OPPORTUNITY

Rails you can actually check

The gap isn't a smarter signal. It's automation whose limits are measured, so real money behind it stops being a leap of faith.

- The income is documented research: index options price in more volatility than actually shows up (Bakshi & Kapadia 2003, Carr & Wu 2009, CBOE's PUT index)
- What's missing is proof the safety layer holds, so Argus replays every guard against real history instead of assuming it works
- That turns a safety claim into a number a buyer can check, not a promise they have to take



THE PREMIUM IS THE GAP BETWEEN THEM

SECTION 02

The Product

What Argus actually does, how it decides, and proof it's trading for real.

WHAT IT ACTUALLY DOES

Insurance on the S&P 500

Sold on a schedule, without a person clicking anything. Every position is a defined-risk put credit spread.

- Sells one put and buys a cheaper, further-out put, at the same time, as a single trade
- That second leg fixes the most Argus can lose the moment the trade is placed
- Both legs always go in together, so there's no window where only one side exists

LEG 1

Sell one put option

LEG 2

Buy a cheaper put, further out

ONE ORDER

Max loss fixed before it exists

WHY THIS BEATS A TYPICAL TRADING BOT

Not your typical trading bot

Most bots let the AI trade freely, and publish one number with no way to check it.
Argus does neither.

	TYPICAL BOT	ARGUS
Reporting	Reports a backtest Sharpe, full stop	Reports the same number, and what it becomes once corrected for every strategy actually tried
AI's power	Can size or place a trade on its own	Can only shrink or cancel a trade a fixed rules engine already built
Safety guards	Assumed to work	Backtested against real historical winning weeks

HOW IT DECIDES

How Argus decides

Fifteen checks have to pass before an AI even sees the trade, and it can only shrink or cancel, never grow it.



Picker

fixed rules engine

Builds the trade proposal



Guards

15 automated checks


Account risk, liquidity, volatility



Reviewer

AI, shrink / veto only

No network call in the clamp
that enforces it



Execution

Alpaca


Multi-leg order, unattended

EXECUTION AND DISCLOSURE

A real fill, fully disclosed

One account proved the system fills correctly. A different account is the one judged, and we say why.

- The fill shown here proves the multi-leg order logic works. It's not an official return: that account was opened two weeks before the hackathon and is disqualified under Alpaca's own rules, regardless of its clean history
- The account actually judged, PA3HWE141FA8, opened Sep 3 after a real signup delay and carries no trades from the scored window. We're stating that plainly rather than leaving it for a judge to find



SECTION 03

The Evaluation

How we stress-test our own claims, including the ones that don't flatter us.

THE TRACK RECORD

A 98.9% win rate

351 wins, 4 losses, across 355 traded weeks from 2016 to 2026. This is the real, replayed result, before any statistical correction.

98.9%

WIN RATE, 351 OF 355
TRADED WEEKS (538
TOTAL, 2016-2026)

2.38

Profit factor: total wins are worth 2.38x total losses

4

Losing weeks out of 355 traded, worst week -4.70%

The next slide asks a harder, different question: not how often trades worked, but whether this is a repeatable skill, or luck from picking this idea out of 31 tried.

SOURCE: PUBLIC/DATA/TRACK_RECORD.JSON (STATS.WIN_RATE, N_WINS, N_LOSSES, PROFIT_FACTOR)

TESTED, NOT JUST CLAIMED

We benchmark our own results

The 98.9% win rate is real. This asks a different question: is it a repeatable skill, or did we get lucky picking this idea out of 31 tried?

DEFLATED SHARPE RATIO

0.20

At N = 31 ideas tried. Not proven yet.

RAW SHARPE, UNCORRECTED

+0.574

Before the multiple-testing
correction

That correction once caught a real mistake in our own simulation, and forced a lower, honest number instead of a flattering one.

SOURCE: WRITEUP.MD PERFORMANCE TABLE; EXPERIMENT_29_SHARPE_AUDIT.MD. REPRODUCE IT YOURSELF: PYTHON -M PIPELINE.FALSEIFY.AUDIT, NO LOGIN, UNDER TWENTY SECONDS.

SAFETY CHECKS ARE TESTED TOO

A guard blocked too much

One of Argus's own guards was quietly blocking 42-50% of good trades, until it was measured and corrected.

- A guard nobody checks against the trades it blocks is a guess wearing a safety label
- Each of the fifteen guards is replayed against real historical winning weeks to measure exactly that

BEFORE ~~42-50%~~ of winning trades blocked

TOO TIGHT


AFTER **0-1.6%** of winning trades blocked

MEASURED, NOT ASSUMED

WHY THIS MATTERS BEYOND ONE TRADE

A reusable honesty engine

Not built for one strategy. Any trading idea can be run through the same gauntlet that caught Argus's own bug.



- `pipeline/falsify/engine.py` runs any trading hypothesis through the same gauntlet, and reports exactly which stage killed it, if any did
- A team building a completely different strategy can point this engine at their own result and get the same honest accounting

WHAT'S NEXT

Two markets, checked first

Before either got real money. One cleared the bar.
One didn't, so it isn't being traded.

The same cheap, honest gate ran before either got real money, and IWM's margin was real enough to keep it on paper.

**QQQ**

CLEARED

IWM

NOT YET

IN SUMMARY

Honest about what's proven

Three things this deck showed, and how to check every one of them yourself.

THE PROBLEM

Options income works. What fails is the risk discipline, manual or automated, that nobody ever tested.

THE PRODUCT

Real trades, fixed rules, and an AI that can only shrink or veto a position, never grow one.

THE EVALUATION

A 98.9% win rate, checked by a self-audit that also caught and fixed our own math before we published it.

Repo · [WRITEUP.md](#) · [Live dashboard](#) · Reproduce it: `python -m pipeline.falsify.audit`

Argus

An AI-assisted trading agent
that sells insurance on the S&P
500.

Repo · WRITEUP.md · Live dashboard

